



Symbol

Security Assessment

27th July, 2020

Prepared For:

Iain Wilson | *NEM Group*
info@nem.group

Dave Hodgson | *NEM Group*
info@nem.group

David Mansell | *NEM Group*
info@nem.group

Prepared By:

Dominik Czarnota | *Trail of Bits*
dominik.czarnota@trailofbits.com

Jim Miller | *Trail of Bits*
jim.miller@trailofbits.com

David Pokora | *Trail of Bits*
david.pokora@trailofbits.com

Changelog:

Jul 27, 2020:	Initial report draft
Aug 7, 2020:	Updated issue #11
Aug 31, 2020:	Added Appendix D. Fuzzing catapult-server
Sep 18, 2020:	Updates made to Appendix D
Sep 24, 2020:	Added Appendix E. Fix Log
Oct 30, 2020:	Review of amendments detailed in Appendix F
Nov 25, 2020:	Review of amendments detailed in Appendix G
Dec 7, 2020:	Added Appendix H and extended executive summary

[Executive Summary](#)

[Project Dashboard](#)

[Engagement Goals](#)

[Coverage](#)

[Recommendations Summary](#)

[Short term](#)

[Long term](#)

[Findings Summary](#)

- [1. Missing compiler mitigations](#)
- [2. Undefined behavior dereferencing std::list.back\(\) on an empty container](#)
- [3. Current ConfigurationBags verification may lead to bugs](#)
- [4. High-entropy RNG does not guarantee high entropy](#)
- [5. Use O_CLOEXEC flag by default when opening files on Linux](#)
- [6. The symbol-cli saves the config file as readable for others](#)
- [7. Maximum packet size of 4GB may lead to denial-of-service attacks](#)
- [8. Lack of overflow checks](#)
- [9. The boost::filesystem::create_directory defaults to 0777 permissions](#)
- [10. Potential padding oracle attack in AesCbcDecrypt](#)
- [11. Incorrect ReceiptType in catapult-rest](#)

[A. Vulnerability Classifications](#)

[B. Compiler Mitigations](#)

[C. Fuzzing catapult-server](#)

[D. Previous security testing report fixes](#)

[E. Fix Log](#)

[Detailed Fix Log](#)

[F. Amendments to Voting Key Structure](#)

[G. Additional Symbol changes review](#)

[H. Fix log for issues from Appendix G](#)

Executive Summary

From June 13 through June 24, 2020, NEM Group engaged Trail of Bits to review the security of the Symbol network's node and REST gateway repositories. Trail of Bits conducted this assessment over the course of four-person weeks with two engineers working from the provided repositories.

In the first week of the assessment, we employed static analysis techniques such as cppcheck, scan-build, CodeQL, and CLion code inspection while gaining a deeper understanding of the codebase through manual review. Preliminary manual review included potential generic classes of issues such as memory corruption, use of cryptography, data validation, auditing/logging, configuration, and more. This led to the discovery of six findings ranging from informational to high severity.

In the final week, we conducted a deeper analysis of the Symbol network, identifying critical components within the system and pursuing testing against those targets. This led to the discovery of several new issues ranging from informational to low severity. A partial review of data flow and chain state validation revealed insufficient data validation issues related to maximum packet sizes ([TOB-SYM-007](#)) and inflationary mosaic supply ([TOB-SYM-008](#)) within the catapult server, and incorrect receipt types defined within the REST gateway ([TOB-SYM-011](#)). Further review of authentication schemes and access controls revealed additional concerns regarding default permissions used with boost filesystem APIs ([TOB-SYM-009](#)).

Finally, a review of cryptography led to the discovery of a potential padding oracle attack ([TOB-SYM-010](#)). Trail of Bits engineers also began pursuing dynamic testing via libFuzzer and various engines multiplexed via [DeepState](#). Unfortunately, when we wished to employ dynamic analysis, the fuzzing targets were often at a level of depth within the system that required additional overhead for setup.

Overall, we believe the Symbol network repositories reviewed show positive consideration for common classes of security vulnerabilities. As a result, low-hanging fruit seemed less prevalent throughout the codebase, and static analysis techniques reflected few true-positive concerns. Trail of Bits believes there may be latent data validation issues deeper within the system that could be exposed through increased dynamic testing. We recommend integrating a robust fuzzing harness such as [AFL](#) to test security properties in-depth, while employing static analysis tools in the CI/CD build pipeline to uncover newly introduced vulnerabilities and strengthen the system's security posture.

Update September 24, 2020: Trail of Bits reviewed fixes implemented for the issues presented in this report, amendments to voting key structure, and additional changes made to certain repositories. See the results from those reviews in [Appendix E: Fix Log](#), [Appendix F: Amendments](#)

to Voting Key Structure, Appendix G: Additional Symbol changes review, and H: Fix log for issues from Appendix G.

Project Dashboard

Application Summary

Name	Symbol
Version	catapult-server commit: f84eb88727 catapult-rest commit: 25d62f2393
Type	C++, JavaScript
Platforms	Windows, Linux

Engagement Summary

Dates	July 13–July 24, 2020
Method	Whitebox
Consultants Engaged	2
Level of Effort	4 person-weeks

Vulnerability Summary

Total High-Severity Issues	1	■
Total Medium-Severity Issues	1	■
Total Low-Severity Issues	3	■ ■ ■
Total Informational-Severity Issues	3	■ ■ ■
Total Undetermined-Severity Issues	3	■ ■ ■
Total	11	■ ■ ■ ■ ■ ■ ■ ■ ■ ■ ■

Category Breakdown

Access Controls	2	■ ■
Configuration	2	■ ■
Cryptography	2	■ ■
Data Validation	2	■ ■ ■
Denial of Service	1	■
Undefined Behavior	2	■ ■
Total	11	■ ■ ■ ■ ■ ■ ■ ■ ■ ■ ■

Engagement Goals

This engagement was scoped to provide a security assessment of the previously mentioned Symbol network components and their communications, namely the catapult-server and catapult-rest repositories.

Specifically, we sought to answer the following questions:

- Are there standard compiler and platform mitigations that could be leveraged to improve the security posture of Symbol nodes?
- Do the provided repositories have code correctness issues that could lead to memory corruption, arbitrary code execution, etc.?
- Are appropriate access controls set for sensitive information stored at rest?
- Do the target components make appropriate use of encryption to secure sensitive information?
- Is the default configuration for the server overly permissive?
- Are there any insufficient data validation issues that may lead to a denial-of-service attack? Could they be leveraged against the consensus model to take over the network?
- Does logging result in sensitive data exposure? Could the logging system be leveraged to perform a resource exhaustion attack?

Coverage

This section highlights some of the analysis coverage we achieved based on our high-level engagement goals.

- Review of currently configured compiler mitigations revealed the need for additional mitigations that would help prevent an attack against the server. ([TOB-SYM-001](#)).
- In analyzing the codebase for generic classes of vulnerabilities such as memory corruption, use-after-frees, arbitrary code execution, and more, we discovered an issue dereferencing a potentially invalid pointer ([TOB-SYM-002](#)) and a lack of overflow checks ([TOB-SYM-008](#)).
- While investigating the use of cryptography in the system, we discovered a potentially cryptographically insecure RNG ([TOB-SYM-004](#)) and a potential padding oracle attack ([TOB-SYM-010](#)). We also found some positive reinforcements, including appropriate wrapping of dependencies like OpenSSL and proper implementation of EdDSA.
- Review of various file operations revealed permission-based concerns ([TOB-SYM-005](#), [TOB-SYM-006](#), [TOB-SYM-009](#)).

- Default server configurations related to authentication were found to be sufficient. For example, by default, only machine-local connections are allowed, and the maximum ban list size is fairly large.
- Based on our review of improper data validation, we are concerned about ConfigurationBags verification ([TOB-SYM-003](#)), maximum packet sizes potentially leading to denial of service ([TOB-SYM-007](#)), and parsing of an incorrect receipt type in the REST gateway ([TOB-SYM-011](#)).
- A partial review of logging did not indicate that sensitive information such as key material could be leaked.

Recommendations Summary

This section aggregates all the recommendations made during the engagement. Short-term recommendations address the immediate causes of issues. Long-term recommendations pertain to the development process and long-term design goals.

Short term

- ❑ **Enable security mitigations for the catapult-server build using the compiler and linker flags described in the table below.** While the compiler used in the [catapult-service-bootstrap](#) build enables some of the mitigations by default, enabling them explicitly ensures the mitigations are there in case the compiler used has different defaults. [TOB-SYM-001](#)
- ❑ **Refactor the method so the `empty()` check ensures at least one item exists in the container before calling `back()` and dereferencing the underlying data.** [TOB-SYM-002](#)
- ❑ **Validate the exact configuration bag size to prevent issues resulting from unexpected size changes.** [TOB-SYM-003](#)
- ❑ **Replace the internal randomness provider for the `HighEntropyRandomGenerator` with one that explicitly provides cryptographically secure randomness, such as those provided by OpenSSL and LibSodium.** [TOB-SYM-004](#)
- ❑ **Add the `O_CLOEXEC` flag to the flags used for opening files in the `RawFile` class by default and drop the flag explicitly when needed.** This will prevent leaking opened file descriptors into child processes if the node spawns child processes. [TOB-SYM-005](#)
- ❑ **Explicitly set the permissions of the `symbol-cli.config.json` file to `00600` in the `symbol-cli` when [creating](#) or [saving](#) the file.** This will prevent users on the same machine from accessing someone else's account if they access the config file and are able to brute-force the password. [TOB-SYM-006](#)
- ❑ **Set lower maximum default packet sizes in the `PacketPayloadBuilder` and `ServerPacketHandlers` classes' constructors.** This will help prevent unwanted denial-of-service scenarios even if an attacker is able to construct big enough payloads to cause them. [TOB-SYM-007](#)
- ❑ **Don't use `boost::filesystem::create_directory` or `create_directories` functions for creating directories in catapult-server.** Instead, use a method that allows you to set permissions of the created directories explicitly and appropriately. This will prevent

insecure configurations that may lead to unexpected behavior and allow access from other system users in less securely configured systems. [TOB-SYM-009](#)

❑ **Include an HMAC alongside the AES-CBC encryption to ensure that ciphertexts are not altered by an adversary and to prevent potential padding oracle attacks.**

[TOB-SYM-010](#)

❑ **Fix the duplicate ReceiptType value in the catapult-rest codebase.** [TOB-SYM-011](#)

Long term

- ❑ **Enable the security mitigations for all binaries used and add scanning for security mitigations into the CI/CD pipeline to ensure that certain options are always enabled.** This will make it harder to exploit potential bugs found in dependencies used by the catapult-server. [TOB-SYM-001](#)
- ❑ **Review the possible return values of such functions to ensure the expected data will always be returned or otherwise handled accordingly.** [TOB-SYM-002](#)
- ❑ **Be mindful of the importance of using libraries that may resolve to implementations that do not satisfy requirements when using a given platform or compiler.** [TOB-SYM-004](#)
- ❑ **Check the permissions of the `symbol-cli.config.json` file when reading it in the `symbol-cli` and warn the user if the file permissions are overly broad.** This will help prevent situations in which the user accidentally sets permissions that are too broad for the file and never changes them back. [TOB-SYM-006](#)
- ❑ **Add overflow checks that would either error out the program or log an error to `BaseValue` operators, `ChainScore` operators, the `MosaicEntrySupplyMixin::increaseSupply` function, and others.** If those checks are unwanted due to performance issues, enable them only during tests and debug builds. This will help catch bugs earlier and prevent unwanted overflows that could cause financial risks. [TOB-SYM-008](#)
- ❑ **Consider clearing plaintext buffers upon any authentication or decryption failure to prevent invalid plaintexts from being used.** [TOB-SYM-010](#)

Findings Summary

#	Title	Type	Severity
1	Missing compiler mitigations	Configuration	Low
2	Undefined behavior dereferencing std::list.back() on an empty container	Undefined Behavior	Undetermined
3	Current ConfigurationBags verification may lead to bugs	Data Validation	Informational
4	High-entropy RNG does not guarantee high entropy	Cryptography	Medium
5	Use O_CLOEXEC flag by default when opening files on Linux	Configuration	Informational
6	The symbol-cli saves the config file as readable for others	Access Controls	High
7	Maximum packet size of 4GB may lead to denial-of-service attacks	Denial of Service	Undetermined
8	Lack of overflow checks	Data Validation	Informational
9	The boost::filesystem::create_directory defaults to 0777 permissions	Access Controls	Low
10	Potential padding oracle attack in AesCbcDecrypt	Cryptography	Undetermined
11	Incorrect ReceiptType in catapult-rest	Undefined Behavior	Low

Figure TOB-SYM-001.1: checksec output. Note that the yellow DSO (dynamic shared object) is ok, as dynamic libraries are always built with PIC (position-independent code), which is the same as PIE (position-independent executable).

Modern compilers support a number of exploit mitigations, including:

- NX (non-executable data)
- PIE (position-independent code for ASLR)
- stack canaries (for buffer overflow detection)
- RELRO (for hardening data sections)
- FORTIFY_SOURCE (for additional buffer overflow detection and format string protection)
- stack clash protection (for detecting when stack pointer clashes with other memory region)
- CFI (control flow integrity)
- SafeStack (for further stack-overflow protection)

By default, compilers do not enable many of these mitigations. For a detailed description of these exploit mitigation technologies, see [Appendix B](#).

Recommendation

Short term, enable security mitigations for the catapult-server build using the compiler and linker flags described in the table below. While the compiler used in the [catapult-service-bootstrap](#) build enables some of the mitigations by default, enabling them explicitly ensures the mitigations are there in case the compiler used has different defaults.

For additional assurance, consider verifying if ASLR is enabled during program startup by checking if the value stored in the `/proc/sys/kernel/randomize_va_space` file is 2, and error out if it is lower.

Long term, enable the security mitigations for all binaries used and add scanning for security mitigations into the CI/CD pipeline to ensure that certain options are always enabled. This will make it harder to exploit potential bugs found in dependencies used by the catapult-server.

GCC flag	What it enables
-z noexecstack	NX bit
-Wl,-z,relro,-z,now	Full RELRO (read-only segments after relocation and disables lazy bindings)
-fstack-protector-all	Adds stack canaries for all the functions. (Note: This might make the program

<p>or (less secure)</p> <p><code>-fstack-protector-strong --param ssp-buffer-size=4</code></p>	<p>slower. To protect only functions that have buffers, use the latter version.)</p>
<p><code>-fPIE -pie</code></p>	<p>PIE (needs ASLR enabled when launching a program).</p>
<p><code>-D_FORTIFY_SOURCE=2 -O2</code></p> <p>or (less secure)</p> <p><code>-D_FORTIFY_SOURCE=1 -O1</code></p>	<p>FORTIFY_SOURCE protections. (Note that it requires an appropriate optimization flag (-O1 or -O2). The latter version (-D_FORTIFY_SOURCE=1 -O1) is less secure as it will enable only compile-time protections, while the former will also add runtime checks.)</p>
<p><code>-fstack-clash-protection</code></p>	<p>Adds checks to functions that may allocate a lot of memory on the stack to ensure the new stack pointer (and stack frame) do not overlap with another memory region such as heap.</p>
<p><code>-fsanitize=cfi -fvisibility=hidden -flto</code></p>	<p>Enables control flow integrity checks that help prevent program control flow hijacking (Clang/LLVM only).</p>
<p><code>-fsanitize=safe-stack</code></p>	<p>Enables SafeStack which splits stack frames of certain functions into the safe stack and the unsafe stack, to make it harder to hijack the program's control flow (Clang/LLVM only).</p>
<p><code>-Wall -Wextra -Wpedantic -Wshadow -Wconversion -Wformat-security</code></p>	<p>Compile-time checks and warnings.</p>

References

- [Debian hardening recommendations](#)
- [GCC man page](#)
- [LD man page \(see -z keywords\)](#)

2. Undefined behavior dereferencing `std::list.back()` on an empty container

Severity: Undetermined

Difficulty: Undetermined

Type: Undefined Behavior

Finding ID: TOB-SYM-002

Target: catapult-server/src/catapult/ionet/NodeInteractionsContainer.cpp

Description

The `NodeInteractionsContainer` has an operation to add a `NodeInteractionBucket` to a `std::list`, but it dereferences `std::list.back()` prior to checking `std::list.empty()`, which may produce undefined behavior.

```
void NodeInteractionsContainer::addInteraction(Timestamp timestamp, const
consumer<NodeInteractionsBucket&>& consumer) {
    auto bucketAge = utils::TimeSpan::FromDifference(timestamp,
m_buckets.back().CreationTime);
    if (m_buckets.empty() || BucketDuration() <= bucketAge)
        m_buckets.push_back(NodeInteractionsBucket(timestamp));

    consumer(m_buckets.back());
}
```

Figure 2.1: Undefined behavior dereferencing `std::list.back()` on an empty container ([catapult-server/src/catapult/ionet/NodeInteractionsContainer.cpp#L71-L77](#)).

The C++ reference noted below states: “Calling this function on an empty container causes undefined behavior.” For example, depending on the platform, this could produce a denial of service by accessing protected/non-existent memory regions.

The reachability of this code path or likelihood of reproduction was not determined, so this issue is marked undetermined severity.

Exploit Scenario

Bob is an operator of a public Symbol network node who expects uninterrupted service due to the perceived nature of Symbol’s security and decentralization. However, due to this bug, Bob’s node encounters a crash. This may also affect other operators like Bob and their ability to perform consensus, which may have dire consequences for network governance.

Recommendation

Short term, refactor the method so the `empty()` check ensures at least one item exists in the container prior to calling `back()` and dereferencing the underlying data.

Long term, review the possible return values of such functions to ensure the expected data will always be returned or otherwise handled accordingly.

References

- [list::back - C++ Reference](#)

3. Current ConfigurationBags verification may lead to bugs

Severity: Informational

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SYM-003

Target: catapult-server/src/catapult/utils/ConfigurationUtils.cpp

Description

The catapult-server uses the VerifyBagSizeLte function to verify the size of the loaded configuration bag. This function only throws an error if the configuration bag is bigger than expected. Therefore, since adding a new parameter will still pass the verification without changing the expected size, any changes to the number of loaded parameters might lead to an unwanted situation.

```
void VerifyBagSizeLte(const ConfigurationBag& bag, size_t expectedSize) {  
    if (bag.size() > expectedSize)  
        CATAPULT_THROW_INVALID_ARGUMENT_1("configuration bag contains too many properties",  
bag.size());  
}
```

Figure TOB-SYM-003.1: The VerifyBagSizeLte function

(catapult-server/blob/v0.9.6.3/src/catapult/utils/ConfigurationUtils.cpp#L38-L41).

Recommendation

Short term, validate the exact configuration bag size to prevent issues resulting from unexpected size changes.

4. High-entropy RNG does not guarantee high entropy

Severity: Medium

Difficulty: Low

Type: Cryptography

Finding ID: TOB-SYM-004

Target: catapult-server/src/catapult/utils/RandomGenerator.cpp

Description

The `HighEntropyRandomGenerator` is used as a source of cryptographically secure randomness. Unfortunately, this is based on `std::random_device`, which is not guaranteed to be high-entropy or cryptographically secure.

Without explicit hardening of standard libraries, different compilers or target platforms may not offer a cryptographically secure source of randomness via `std::random_device`. The provider may even be deterministic.

Exploit Scenario

Bob is a node operator who wishes to run Symbol network nodes on a given platform. An attacker, Eve, knows that `std::random_device` is not cryptographically secure on the platform Bob's nodes are running on, so she may be able to deduce Bob's private keys more easily, leaving Bob's funds at risk.

Recommendation

Short term, replace the internal randomness provider for the `HighEntropyRandomGenerator` with one that explicitly provides cryptographically secure randomness, such as those provided by OpenSSL and LibSodium.

Long term, be mindful of the importance of using libraries that may resolve to implementations that do not satisfy requirements when using a given platform or compiler.

References

- [Is std::random_device cryptographic\[ally\] secure?](#)
- [Everything You Never Wanted to Know about C++'s random_device: It Might Actually Be Deterministic](#)

5. Use O_CLOEXEC flag by default when opening files on Linux

Severity: Informational

Difficulty: Undetermined

Type: Configuration

Finding ID: TOB-SYM-005

Target: catapult-server/src/catapult/io/RawFile.cpp

Description

The RawFile wrapper used for opening files does not use the O_CLOEXEC flag by default. Enabling this flag allows an opened file to close when a process calls execve and thereby prevents unexpected leaks of file descriptors to child processes.

```
namespace catapult { namespace io {
    namespace {
        // (...)
#ifdef _MSC_VER
        // (...)
#else
        // (...)
        constexpr auto Flag_Read_Only = O_RDONLY;
        constexpr auto Flag_Read_Write = O_RDWR;
        constexpr auto New_File_Create_Truncate_Flags = O_CREAT | O_TRUNC;
        constexpr auto New_File_Create_Flags = O_CREAT;

        FileOperationResult<int> nemOpen(int& fd, const char* name, OpenMode mode, LockMode
lockMode) {
            int flags = mode == OpenMode::Read_Only ? Flag_Read_Only : Flag_Read_Write;
            int createFlag = mode == OpenMode::Read_Write
                ? New_File_Create_Truncate_Flags
                : (mode == OpenMode::Read_Append ? New_File_Create_Flags : 0);
            int lockingFlags = LockMode::File == lockMode
                ? ((flags & Flag_Read_Write) ? File_Locking_Exclusive :
File_Locking_Shared_Read)
                : File_Locking_None;

            return open(fd, name, File_Binary_Flag | flags | createFlag, lockingFlags,
New_File_Permissions);
        }
    }
}
```

Figure TOB-SYM-005.1: File flags and their usage

([catapult-server/blob/v0.9.6.3/src/catapult/io/RawFile.cpp#L134-L137](#) and [#L230-L240](#)).

Recommendation

Short term, add the O_CLOEXEC flag to the flags used for opening files in the RawFile class by default and drop the flag explicitly when needed. This will prevent leaking opened file descriptors into child processes if the node spawns child processes.

6. The `symbol-cli` saves the config file as readable for others

Severity: High

Type: Access Controls

Target: `symbol-cli`

Difficulty: Low

Finding ID: TOB-SYM-006

Description

Importing a profile into `symbol-cli` creates the `symbol-cli.config.json` file. This file is created with the `fs.writeFileSync` function (Figure TOB-SYM-006.1), which [by default uses the 0666 file permissions](#). This allows all users to read the created config file.

```
/**
 * Save profiles from JSON.
 * @param {JSON} profiles
 */
private saveProfiles(profiles: ProfileRecord) {
    fs.writeFileSync(this.filePath, JSON.stringify(profiles), 'utf-8');
}
```

Figure TOB-SYM-006.1: The `saveProfiles` function that saves the `symbol-cli.config.json` file with permissions that are too broad ([symbol-cli/src/respositories/profile.repository.ts#L203-L209](#)).

Exploit Scenario

Alice imports her account into `symbol-cli`. Eve, who has access to Alice's home directory, copies her `symbol-cli.config.json` file and brute-forces their password. Eve then steals Alice's mosaics.

Recommendation

Short term, explicitly set the permissions of the `symbol-cli.config.json` file to `0600` in the `symbol-cli` when [creating](#) or [saving](#) the file. This will prevent users on the same machine from accessing someone else's account if they access the config file and are able to brute-force the password.

Long term, check the permissions of the `symbol-cli.config.json` file when reading it in the `symbol-cli` and warn the user if the file permissions are too broad. This will help prevent situations in which the user accidentally sets permissions that are too broad for the file and never changes them back.

7. Maximum packet size of 4GB may lead to denial-of-service attacks

Severity: Undetermined

Difficulty: High

Type: Denial of Service

Finding ID: TOB-SYM-007

Target: catapult-server/../../PacketPayloadBuilder.h and PacketHandlers.h

Description

The PacketPayloadBuilder and ServerPacketHandlers classes' constructors set the maxPacketDataSize argument's default value to a maximum 32-bit unsigned integer value, which is more than 4GB (Figures TOB-SYM-007.1-2). Allowing such big packet sizes may lead to denial-of-service attacks in which an attacker would fill the node's memory with huge packets.

```
/// Packet payload builder for creating payloads composed of heterogeneous data.
class PacketPayloadBuilder {
public:
    /// Creates builder for a packet with the specified \a type.
    explicit PacketPayloadBuilder(PacketType type) : PacketPayloadBuilder(type,
std::numeric_limits<uint32_t>::max())
    {}

    /// Creates builder for a packet with the specified \a type and max packet data size (\a
maxPacketDataSize).
    PacketPayloadBuilder(PacketType type, uint32_t maxPacketDataSize)
        : m_maxPacketDataSize(maxPacketDataSize)
        , m_payload(type)
        , m_hasError(false)
    {}
}
```

Figure TOB-SYM-007.1: Setting the default maxPacketDataSize in PacketPayloadBuilder class' constructor

([catapult-server/src/catapult/ionet/PacketPayloadBuilder.h#L28-L40](#)).

```
/// Collection of packet handlers where there is at most one handler per packet type.
class ServerPacketHandlers {
// (...)
public:
    /// Creates packet handlers with a max packet data size (\a maxPacketDataSize).
    explicit ServerPacketHandlers(uint32_t maxPacketDataSize =
std::numeric_limits<uint32_t>::max());
}
```

Figure TOB-SYM-007.2: Setting the default maxPacketDataSize in ServerPacketHandlers class' constructor ([catapult-server/src/catapult/ionet/PacketHandlers.h#L68-L79](#)).

Recommendation

Short term, set lower maximum default packet sizes in the PacketPayloadBuilder and ServerPacketHandlers classes' constructors. This will help prevent unwanted denial-of-service scenarios even if an attacker is able to construct big enough payloads to cause them.

8. Lack of overflow checks

Severity: Informational
Type: Data Validation
Target: catapult-server

Difficulty: Undetermined
Finding ID: TOB-SYM-008

Description

The catapult-server provides wrappers for various value types and simplifies their usage with arithmetic operator overloads. However, many of those operations are not checked against integer overflows and may lead to unexpected behavior if an overflow does occur.

While we did not find any cases where such an overflow could occur, the BaseValue operators (Figure TOB-SYM-008.1), ChainScore operators (Figure TOB-SYM-008.2), MosaicEntrySupplyMixin::increaseSupply function (Figure TOB-SYM-008.3), and others could use the CheckedAdd utility function (Figure TOB-SYM-008.4) and either error out or log an error if an overflow occurs.

```
constexpr BaseValue operator+(BaseValue rhs) const {  
    return BaseValue(this->unwrap() + rhs.unwrap());  
}  
  
/// Subtracts \a rhs from this value and returns a new value.  
constexpr BaseValue operator-(BaseValue rhs) const {  
    return BaseValue(this->unwrap() - rhs.unwrap());  
}
```

Figure TOB-SYM-008.1: The BaseValue operators
([catapult-server/src/catapult/utils/BaseValue.h#L109-L116](#)).

```
/// Adds \a rhs to this chain score.  
ChainScore& operator+=(const ChainScore& rhs) {  
    m_score += rhs.m_score;  
    return *this;  
}  
  
/// Subtracts \a rhs from this chain score.  
ChainScore& operator-=(const ChainScore& rhs) {  
    m_score -= rhs.m_score;  
    return *this;  
}
```

Figure TOB-SYM-008.2: The ChainScore operators
([catapult-server/src/catapult/model/ChainScore.h#L65-L75](#)).

```
void MosaicEntrySupplyMixin::increaseSupply(Amount delta) {  
    m_supply = m_supply + delta;  
}
```

Figure TOB-SYM-008.3: The MosaicEntrySupplyMixin::increaseSupply function
([catapult-server/plugins/txes/mosaic/src/state/MosaicEntry.cpp#L31-L40](#)).

```
/// Adds \a delta to \a value if and only if there is no overflow.
template<typename T>
bool CheckedAdd(T& value, T delta) {
    if (value > std::numeric_limits<T>::max() - delta)
        return false;

    value += delta;
    return true;
}
```

Figure TOB-SYM-008.4: The CheckedAdd function
([catapult-server/src/catapult/utils/IntegerMath.h#L28-L36](#)).

Recommendation

Long term, add overflow checks that would either error out the program or log an error to BaseValue operators, ChainScore operators, the MosaicEntrySupplyMixin::increaseSupply function, and others. If those checks are unwanted due to performance issues, enable them only during tests and debug builds. This will help catch bugs earlier and prevent unwanted overflows that could cause financial risks.

9. The `boost::filesystem::create_directory` defaults to `0777` permissions

Severity: Low

Type: Access Controls

Target: catapult-server

Difficulty: High

Finding ID: TOB-SYM-009

Description

The [boost::filesystem::create_directory function](#) used across the catapult-server codebase creates a directory through the [mkdir syscall](#) passing in `0777` permissions. While this is usually further limited by the default [umask setting](#) set on a given Linux system, this might allow incorrect permissions for the created directories, and an attacker who has an account on the same system would be able to create files and directories in there.

Exploit Scenario

Alice sets up her machine with a `umask` setting that's too broad and hosts a Symbol node. Eve, who has an account on the same machine as Alice, accesses directories created by Alice's Symbol node and alters the behavior of that node.

Recommendation

Short term, don't use `boost::filesystem::create_directory` or `create_directories` functions for creating directories in catapult-server. Instead, use a method that allows you to set permissions of the created directories explicitly and appropriately. This will prevent insecure configurations that may lead to unexpected behavior and allow access from other system users in less securely configured systems.

10. Potential padding oracle attack in AesCbcDecrypt

Severity: Undetermined

Difficulty: Undetermined

Type: Cryptography

Finding ID: TOB-SYM-010

Target: catapult-server/src/catapult/crypto/AesCbcDecrypt.cpp

Description

The catapult-server uses an elliptic curve integrated encryption scheme (ECIES) in its crypto library. Specifically, Ed25519 is used to derive a shared secret, which acts as an input into a key derivation function, which is then used as a symmetric key for AES-CBC. However, AES-CBC is paired with PKCS #7 padding and is not authenticated with a message authentication code (MAC), which could make a [padding oracle attack](#) possible.

```
bool TryAesCbcDecrypt(const SharedKey& key, const RawBuffer& input, std::vector<uint8_t>&
output) {
    AesInitializationVector initializationVector;
    if (input.Size < initializationVector.size())
        return false;

    output.resize(input.Size - initializationVector.size());
    if (0 != output.size() % Aes_Pkcs7_Padding_Size)
        return false;

    std::memcpy(initializationVector.data(), input.pData, initializationVector.size());

    auto outputSize = static_cast<int>(output.size());

    OpensslCipherContext cipherContext;
    cipherContext.dispatch(EVP_DecryptInit_ex, EVP_aes_256_cbc(), nullptr, key.data(),
initializationVector.data());
    cipherContext.dispatch(EVP_DecryptUpdate, output.data(), &outputSize, input.pData +
initializationVector.size(), outputSize);
    if (!cipherContext.tryDispatch(EVP_DecryptFinal_ex, output.data() + outputSize,
&outputSize))
        return false;

    // drop PKCS#7 padding
    if (!DropPadding(output))
        return false;

    return true;
}
```

Figure TOB-SYM-010.1: AES-CBC used without a MAC
([catapult-server/src/catapult/crypto/AesCbcDecrypt.cpp#L58-L82](#)).

The catapult-server uses ECIES inside of the registerServices function for its HarvestingServiceRegistrar. If an attacker were able to repeatedly alter and transmit ciphertexts, this function could be turned into a padding oracle, allowing an attacker to recover the underlying plaintexts.

It is unclear exactly how much access an adversary would have to the ciphertexts, and thus it's unclear if this issue is currently exploitable.

Exploit Scenario

An attacker, Eve, is able to alter and send multiple ciphertexts to be decrypted inside of the `registerServices` function. Eve observes enough decryption failures to learn the underlying plaintexts of those ciphertexts.

Recommendation

Short term, include an HMAC alongside the AES-CBC encryption to ensure that ciphertexts are not altered by an adversary and to prevent potential padding oracle attacks.

Long term, consider clearing plaintext buffers upon any authentication or decryption failure to prevent invalid plaintexts from being used.

11. Incorrect ReceiptType in catapult-rest

Severity: Low

Difficulty: Undetermined

Type: Undefined Behavior

Finding ID: TOB-SYM-011

Target: catapult-rest/catapult-sdk/src/plugins/receipts.js

Description

The ReceiptType defined in catapult-rest contains a duplicate value (Figure TOB-SYM-011.1) which is inconsistent with the ReceiptType defined in catapult-server (Figure TOB-SYM-011.2).

```
const ReceiptType = {
  1: 'receipts.balanceTransfer',
  2: 'receipts.balanceChange',
  3: 'receipts.balanceChange',
  4: 'receipts.artifactExpiry',
  5: 'receipts.inflation'
};
```

Figure TOB-SYM-011.1: ReceiptType in catapult-rest
([catapult-rest/catapult-sdk/src/plugins/receipts.js#L24-L30](#)).

```
/// Enumeration of basic receipt types.
/// \note BasicReceiptType is used as highest nibble of receipt type.
enum class BasicReceiptType : uint8_t {
    /// Some other receipt type.
    Other = 0x0,

    /// Balance transfer.
    BalanceTransfer = 0x1,

    /// Balance credit.
    BalanceCredit = 0x2,

    /// Balance debit.
    BalanceDebit = 0x3,

    /// Artifact expiry receipt.
    ArtifactExpiry = 0x4,

    /// Inflation.
    Inflation = 0x5,

    /// Aggregate receipt.
    Aggregate = 0xE,

    /// Alias resolution.
    AliasResolution = 0xF
};
```

Figure TOB-SYM-011.2: ReceiptType in catapult-server
([catapult-server/src/catapult/model/ReceiptType.h#L29-L55](#)).

Recommendation

Short term, fix the duplicate `ReceiptType` value in the `catapult-rest` codebase.

Trail of Bits discussed this issue with the Symbol team and agreed that this issue does not require the proposed change. The `BalanceCredit` and `BalanceDebit` receipt types are included in the `balanceChange` type and don't need separate representations. Due to that, we only recommend documenting this in the code.

A. Vulnerability Classifications

Vulnerability Classes	
Class	Description
Access Controls	Related to authorization of users and assessment of rights
Auditing and Logging	Related to auditing of actions or logging of problems
Authentication	Related to the identification of users
Configuration	Related to security configurations of servers, devices or software
Cryptography	Related to protecting the privacy or integrity of data
Data Exposure	Related to unintended exposure of sensitive information
Data Validation	Related to improper reliance on the structure or values of data
Denial of Service	Related to causing system failure
Error Reporting	Related to the reporting of error conditions in a secure fashion
Patching	Related to keeping software up to date
Session Management	Related to the identification of authenticated users
Timing	Related to race conditions, locking or order of operations
Undefined Behavior	Related to undefined behavior triggered by the program

Severity Categories	
Severity	Description
Informational	The issue does not pose an immediate risk, but is relevant to security best practices or Defense in Depth
Undetermined	The extent of the risk was not determined during this engagement
Low	The risk is relatively small or is not a risk the customer has indicated is important
Medium	Individual user's information is at risk, exploitation would be bad for client's reputation, moderate financial impact, possible legal implications for client

High	Large numbers of users, very bad for client's reputation, or serious legal or financial implications
------	--

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploit was not determined during this engagement
Low	Commonly exploited, public tools exist or can be scripted that exploit this flaw
Medium	Attackers must write an exploit, or need an in-depth knowledge of a complex system
High	The attacker must have privileged insider access to the system, may need to know extremely complex technical details or must discover other weaknesses in order to exploit this issue

B. Compiler Mitigations

NX makes the data sections (including the stack and heap) of the program non-executable. This makes it more difficult for an attacker to execute shellcode. Attackers normally use return-oriented programming (ROP) to bypass NX. Forcing attackers to use ROP makes exploits less reliable across different builds of a program. This mitigation is enabled by default.

Stack canaries (also known as stack cookies) make buffer overflow vulnerabilities more difficult to exploit. A stack canary is a global, randomly-generated value that is copied to the stack between the stack variables and stack metadata in a function's prologue. When a function returns, the canary on the stack is checked against the global value. The program exits if there's a mismatch. This makes it more difficult for an exploit to overwrite the return address on the stack. Depending on the circumstances, attackers may bypass this mitigation by leaking the cookie with a separate information leak vulnerability or by brute-forcing the cookie byte-by-byte.

ASLR (address-space layout randomization) randomizes where each section of the program is placed in memory. This makes it more difficult for an attacker to write reliable exploits, primarily by making it more difficult to jump to ROP gadgets. ASLR requires cooperation from both the system and the compiler. In order to support ASLR fully, a program must be compiled as a position-independent executable (PIE). Most of the Linux distributions have ASLR enabled. This can be checked by reading the value stored in the `/proc/sys/kernel/randomize_va_space` file: 0 means that ASLR is disabled, 1 means it is partially enabled (fewer bits of the addresses are randomized), and 2 means it is fully enabled. This file is writable, and an admin can disable or enable this mitigation. ASLR may be bypassed if an attacker has an information leak in the program.

RELRO (relocations read-only) is a mitigation technique to harden the data sections of an ELF process. It has three modes of operation: disabled, partial, and full. When a program uses a function from a dynamically loaded library, this function address is stored in the GOT.PLT section (Global Offset Table for Procedure Linkage Table). When RELRO is disabled, the function addresses in GOT.PLT point to a dynamic resolver function that resolves the given function address when it is called for the first time. In this case, the memory where the address is stored is both readable and writable. Because of that, an attacker who has control over the process control flow can change the entry of a given function in GOT.PLT to point to any other executable address. For example, they can change the `puts` function's GOT.PLT entry to point to a system function. Then, if the program calls `puts("/bin/sh")`, a `system("/bin/sh")` would be called instead. When RELRO is fully enabled, the dynamic resolver resolves all of the addresses on program startup and changes the permissions of data sections (and therefore GOT.PLT) to read-only.

FORTIFY_SOURCE is a glibc-specific feature that enables a series of mitigations primarily aimed at preventing buffer overflows. With a FORTIFY_SOURCE level of 1, glibc will add compile-time warnings when potentially unsafe calls to common libc functions (e.g., memcpy and strcpy) are made. With a FORTIFY_SOURCE level of 2, glibc will add more stringent run-time checks to these functions. Additionally, glibc will enable a number of lesser-known mitigations. For example, it will disallow the use of a %n format specifier in format strings that aren't located in read-only memory pages. This prevents overwriting data (and gaining code execution) with format string vulnerabilities.

Stack clash protection mitigates a "stack clash vulnerability," where a program's stack memory region grows so much that it overlaps with another memory region. This bug makes the program confuse two different memory addresses (stack and, e.g., heap) so that some of their data overlap, leading to denial of service or control flow hijacking. The stack clash protection mitigation works by adding explicit memory probing to functions that allocate a lot of stack memory, so that the function's stack allocation will never make the stack pointer jump over the stack memory guard page, which is located before the stack.

Control Flow Integrity (CFI)

(<https://clang.llvm.org/docs/ControlFlowIntegrity.html>) mitigates various control-flow hijack attempts and so makes it harder to exploit vulnerabilities like use-after-free or use exploitation techniques such as ROP. This mitigation is currently implemented only in Clang/LLVM.

SafeStack (<https://clang.llvm.org/docs/SafeStack.html>) makes it harder to exploit stack-based buffer overflows as it separates the program stack into the safe stack (which holds saved registers and return addresses) and unsafe stack (which stores everything else). This mitigation is Clang/LLVM only.

C. Fuzzing catapult-server

During the assessment, Trail of Bits attempted to integrate [libFuzzer](#), an in-process, coverage-guided, evolutionary fuzzing engine integrated into Clang. However, due to build issues, most likely related to compiling some libraries with GCC and the final executable with Clang (Figure C.1), [DeepState](#) was integrated instead, and its dumb fuzzing engine (the `--fuzz` flag) was used. DeepState provides an interface for symbolic execution and fuzzing engines through a unit-test-like structure.

Figures C.2-5 show the developed cmake files and the fuzzing harness that performs a round-robin serialization and deserialization of a TransactionInfo structure. The harness needs to be improved to account for finding more paths in the program; currently, with dumb fuzzing, it usually don't pass through [the loop that extracts addresses and crashes there instead](#) due to insufficient buffer length. Figure C.6 shows a traceback from such a crash. Alternatively, instead of refactoring the harness to find more paths, [AFL](#) could be used to find new paths and more interesting crashes as it de-duplicates non-unique crashes.

```
root@7dbccd073699:/host/build# ninja tests.fuzz.initial
[1/1] Linking CXX executable bin/tests.fuzz.initial
FAILED: bin/tests.fuzz.initial
: && /usr/bin/clang++-9 -stdlib=libc++ -Weverything -Werror
-Wno-c++98-compat -Wno-c++98-compat-pedantic
-Wno-disabled-macro-expansion -Wno-padded -Wno-switch-enum
-Wno-weak-vtables -fvisibility=hidden -fsanitize=fuzzer -Wl,--disable-new-dtags
src/catapult/version/nix/CMakeFiles/catapult.version.nix.dir/what_version.cpp.o -o
bin/tests.fuzz.initial
-Wl,-rpath,"$ORIGIN/./deps:$ORIGIN/./lib:/root/boost-build-1.71.0/lib"
/root/boost-build-1.71.0/lib/libboost_atomic.so
/root/boost-build-1.71.0/lib/libboost_system.so
/root/boost-build-1.71.0/lib/libboost_date_time.so
/root/boost-build-1.71.0/lib/libboost_regex.so
/root/boost-build-1.71.0/lib/libboost_timer.so
/root/boost-build-1.71.0/lib/libboost_chrono.so
/root/boost-build-1.71.0/lib/libboost_log.so
/root/boost-build-1.71.0/lib/libboost_thread.so -lpthread
/root/boost-build-1.71.0/lib/libboost_filesystem.so
/root/boost-build-1.71.0/lib/libboost_program_options.so
/root/boost-build-1.71.0/lib/libboost_log_setup.so && :
/usr/bin/ld:
/usr/lib/llvm-9/lib/clang/9.0.0/lib/linux/libclang_rt.fuzzer-x86_64.a(FuzzerDataFlowTrace.cp
p.o): undefined reference to symbol
'_ZNSt14basic_ifstreamIcSt11char_traitsIcEEC1ERKNSt7__cxx1112basic_stringIcS1_SaIcEEEESt13_Io
s_Openmode@@GLIBCXX_3.4.21'
//usr/lib/x86_64-linux-gnu/libstdc++.so.6: error adding symbols: DSO missing from command
line
clang: error: linker command failed with exit code 1 (use -v to see invocation)
ninja: build stopped: subcommand failed.
```

Figure C.1: Build issues when compiling an example libfuzzer fuzzing harness.


```

#include <stdint>
#include <deepstate/DeepState.hpp>

#include "catapult/io/TransactionInfoSerializer.h"
#include "tests/test/core/mocks/MockMemoryStream.h"
#include "catapult/model/Address.h"

using namespace deepstate;
using namespace catapult;

TEST(TInfo, RoundTripSerialization) {
    constexpr const size_t SIZE = 2*sizeof(model::TransactionInfo);
    uint8_t* raw_data = reinterpret_cast<uint8_t*>(DeepState_Malloc(SIZE));
    // Prepare stream for serialization
    std::vector<uint8_t> data{raw_data, raw_data+SIZE};

    mocks::MockMemoryStream stream(data);

    // Deserialize stream into TransactionInfo
    model::TransactionInfo tinfo;
    ReadTransactionInfo(stream, tinfo);

    // Prepare stream for serialization output
    std::vector<uint8_t> buffer;
    mocks::MockMemoryStream stream2(buffer);

    // Serialize tinfo into bytes
    WriteTransactionInfo(tinfo, stream2);

    ASSERT_EQ(buffer.size(), SIZE) << "buffer size is " << buffer.size() << " expected: " <<
    SIZE;
    ASSERT_EQ(memcmp(&data[0], &buffer[0], SIZE), 0) << "Data not same";
}

```

Figure C.2: An example fuzzing harness (tests/deepstate/initial/InitialTests.cpp).

```

function(catapult_deepstate_executable TARGET_NAME)
    find_package(GTest REQUIRED)
    catapult_executable(${TARGET_NAME} ${ARGN})
    add_test(NAME ${TARGET_NAME} WORKING_DIRECTORY ${CMAKE_BINARY_DIR} COMMAND
    ${TARGET_NAME})
endfunction()

function(catapult_deepstate_executable_target TARGET_NAME TEST_DEPENDENCY_NAME)
    catapult_deepstate_executable(${TARGET_NAME} ${ARGN})

    catapult_set_test_compiler_options()
    # Prevent compilation error (see https://github.com/trailofbits/deepstate/issues/357)
    set(CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} -Wno-error" PARENT_SCOPE)

    # Uncomment for compiling with DeepState and libfuzzer; consider also -fsanitize=address
    # target_link_libraries(${TARGET_NAME} -fsanitize=fuzzer -ldeepstate -ldeepstate_LF
    ${TEST_DEPENDENCY_NAME})
    target_link_libraries(${TARGET_NAME} -ldeepstate ${TEST_DEPENDENCY_NAME})
    catapult_target(${TARGET_NAME})
endfunction()

```

Figure C.3: Addition to CMakeGlobalSettings.cmake to use DeepState.

```

cmake_minimum_required(VERSION 3.14)
project(tests.deepstate)

add_subdirectory(initial)

```

Figure C.4: The tests/deepstate/CMakeLists.txt file.

```

cmake_minimum_required(VERSION 3.14)

catapult_deepstate_executable_target(tests.deepstate.initial tests.catapult.test.core)

```

Figure C.5: The tests/deepstate/initial/CMakeLists.txt file.

```

(gdb) bt
#0 boost::log::v2_mt_posix::basic_record_ostream<char>::get_record (this=0x669fd8)
   at /root/boost-build-1.71.0/include/boost/log/sources/record_ostream.hpp:161
#1 0x000000000040fc94 in
boost::log::v2_mt_posix::aux::record_pump<catapult::utils::log::catapult_logger>::~~record_pu
mp (this=0x7fffffffe10)
   at /root/boost-build-1.71.0/include/boost/log/sources/record_ostream.hpp:529
#2 0x00000000004102e9 in catapult::io::BufferInputStreamAdapter<std::__1::vector<unsigned
char, std::__1::allocator<unsigned char> > >::read (
   this=0x7fffffffe2c8, buffer=...) at ../src/catapult/io/BufferInputStreamAdapter.h:59
#3 0x0000000000419fa3 in catapult::io::ReadTransactionInfo (inputStream=...,
transactionInfo=...)
   at ../src/catapult/io/TransactionInfoSerializer.cpp:51
#4 0x000000000040590d in DeepState_Test_TInfo_RoundTripSerialization () at
../tests/deepstate/initial/InitialTests.cpp:28
#5 0x0000000000405879 in DeepState_Run_TInfo_RoundTripSerialization () at
../tests/deepstate/initial/InitialTests.cpp:12
#6 0x000000000040bae3 in DeepState_RunTestNoFork.isra.0 () at
../src/catapult/io/BufferInputStreamAdapter.h:31
#7 0x000000000040bc1c in DeepState_ForkAndRunTest () at
../src/catapult/io/BufferInputStreamAdapter.h:31
#8 0x000000000040bd71 in DeepState_RunSavedTestCase () at
../src/catapult/io/BufferInputStreamAdapter.h:31
#9 0x00000000004053e6 in main ()

```

Figure C.5: Traceback from a crash from running the fuzzing harness from Figure C.2.

D. Previous security testing report fixes

As a secondary goal of our engagement, Trail of Bits was asked to assess remediations in response to the issues highlighted in a report completed by a previous security testing agency.

Due to time constraints, we were unable to test the effectiveness of all the fixes applied by NEM Group. Here are the fixes we had sufficient time to review:

- **#2: Adhoc key derivation scheme allows for the selection of weak keys:** This issue is resolved since HKDF key derivation is now used in `SharedKey.cpp`.
- **#4: Bulk signature verification uses insufficiently random coefficients:** This issue should now be resolved because bulk verification no longer happens at the given target.
- **#5: Custom implementation of Ed25519-SHA3-512:** This is no longer relevant since the Ed25519 implementation included seems to be a standardized one with no custom SHA3-512 implementation. There are notes indicating the move to Ed25519 with standard SHA512 derivation.
- **#8: Protocol content is not confidential:** This concern should be remediated by [this commit](#), which implemented TLS 1.3. An inspection of catapult-service-bootstrap traffic (2dff3fb) does not reveal plaintext data.

Additionally, Trail of Bits engineers discussed the design of a finality protocol with NEM Group. We did not find any issues in the agreed-upon finality proposal, and the finality design is now in Release Code, which should resolve the following issues:

- **#20: Block finality is only (seemingly) guaranteed by the maximum rollback limit:** Concerns regarding this issue should be inherently remediated with the finality roadmap specified by Trail of Bits engineers.
- **#22: Nothing inhibits harvesting on multiple chains (Nothing at Stake):** A finality protocol that maintains a single, finalized chain is being designed and implemented, and does not allow parties to vote for multiple chains in each round. Additionally, they can potentially penalize parties if they do give multiple votes.

E. Fix Log

NEM Group addressed issues TOB-SYM-001 through TOB-SYM-011 in their codebase as a result of our assessment. Each of the fixes was verified by the audit team except [TOB-SYM-008](#), which was partially verified.

ID	Title	Severity	Status
01	Missing compiler mitigations	Low	Fixed
02	Undefined behavior dereferencing <code>std::list.back()</code> on an empty container	Undetermined	Fixed
03	Current ConfigurationBags verification may lead to bugs	Informational	Fixed
04	High-entropy RNG does not guarantee high entropy	Medium	Fixed
05	Use <code>O_CLOEXEC</code> flag by default when opening files on Linux	Informational	Fixed
06	The <code>symbol-cli</code> saves the config file as readable for others	High	Fixed
07	Maximum packet size of 4GB may lead to denial-of-service attacks	Undetermined	Fixed
08	Lack of overflow checks	Informational	Undetermined
09	The <code>boost::filesystem::create_directory</code> defaults to <code>0777</code> permissions	Low	Fixed
10	Potential padding oracle attack in <code>AesCbcDecrypt</code>	Undetermined	Fixed
11	Incorrect <code>ReceiptType</code> in <code>catapult-rest</code>	Low	Unnecessary

Detailed Fix Log

Finding 1: Missing compiler mitigations

Fixed. The appropriate compiler mitigations have been [added to CMakeGlobalSettings.cmake](#) for release builds.

Finding 2: Undefined behavior dereferencing `std::list.back()` on an empty container

Fixed. NEM Group has [added a check](#) to verify if the bucket is empty prior to dereferencing via the `back()` call. If the bucket is empty, a new item will be added, ensuring it is non-empty before dereferencing.

Finding 3: Current ConfigurationBags verification may lead to bugs

Fixed. An [explicit size check was added](#) to ensure the expected size is strictly equal to the provided bag size.

Finding 4: High-entropy RNG does not guarantee high entropy

Fixed. The introduction of the [SecureRandomGenerator](#) alleviates this issue as it uses OpenSSL's random provider under the hood.

Finding 5: Use `O_CLOEXEC` flag by default when opening files on Linux

Fixed. The `O_CLOEXEC` flag has been added to a [definition](#), which is used [during the open\(\) operation](#).

Finding 6: The `symbol-cli` saves the config file as readable for others

Fixed. The affected file has been updated to set appropriate permissions after writing the file. This includes appropriate error handling if the permissions cannot be set, i.e., if the file is removed between writing and setting permissions.

Finding 7: Maximum packet size of 4GB may lead to denial-of-service attacks

Fixed. The maximum packet size has been changed to 100MB. This has not been evaluated in practice but it should reduce DoS attack severity.

Finding 8: Lack of overflow checks

Undetermined. Appropriate overflow checks have been implemented for the second and third cases within the issue description; however, the first case that deals with the `utils::BaseValue` class was repurposed to support many value types. Our fix review has not determined whether every use of this object is secure.

Finding 9: The `boost::filesystem::create_directory` defaults to `0777` permissions

Fixed. This has been fixed by [implementing a CreateDirectory wrapper function](#) that calls the `mkdir` function explicitly passing the `0700` permissions.

Finding 10: Potential padding oracle attack in AesCbcDecrypt

Fixed. [AesCbcDecrypt.cpp](#) has been replaced with [AesDecrypt.cpp](#), which makes use of AES-GCM, which inherently includes a message authentication code (MAC).

Finding 11: Incorrect ReceiptType in catapult-rest

Not fixed, because repair of the issue is unnecessary. These naming conventions may be ambiguous, but do not introduce a security vulnerability. Please refer to the recommendations section of the issue.

F. Amendments to Voting Key Structure

Between October 28 and October 29th, 2020, Trail of Bits reviewed amendments made to the voting key structure. Engineers performed this review working from commit hash [b7b780c](#). At a high level, changes included a transition from the three-level voting key structure to a two-level voting key structure.

These changes align with changes to the specification derived from discussions between NEM Group and Trail of Bits (as mentioned briefly in [Appendix D](#)). The changes largely simplified this portion of the finality process by removing the use of on-the-fly keys and instead opting to sign everything in a given epoch with a specific epoch-tied key. Changes to this process included removing `batchId` and `dilution` variables, and instead using the `keyId` derived from epoch as an alternative. As a result, the “top signature” level was removed, and the focus shifted to verification in root and bottom signatures.

Our review confirmed the implementation matched the agreed-upon specification and ensured general code correctness throughout. Trail of Bits did not identify any issues during this amendment review.

G. Additional Symbol changes review

Between November 23 and November 25th, 2020, Trail of Bits reviewed changes made to the Symbol repositories. An engineer performed this review working from the versions specified by NEMtech, listed in the table below. A non-exhaustive list of updates reviewed included block changes related to importance, as well as the addition of various endpoints and epoch adjustment code. Additionally, we ran a set of Symbol nodes based on the [symbol-bootstrap](#) setup with Docker in order to confirm that nodes properly validate other nodes they connect to.

Repository	Changes reviewed from	Audited version
catapult-server (main branch)	c6f2ffd3a801	8add85f1bd7
catapult-rest (dev branch)	5c6ec9586a5	ed20bcc7cc4
symbol-sdk-typescript-javascript (dev branch)	8916ad9486	2764846fe9f
symbol-sdk-java (dev branch)	2274c70718	60c49f2118c

During the review, we identified five bugs or code quality issues. We describe those findings below, along with recommendations on how to resolve them.

1. The endpoint.Host and metadata.Name are truncated to a length of 255 (through the [GetPackedSize](#) function and a later memcpy), during Node objects serialization in the [PackNode](#) function. In practice, this issue does not expose much risk, as a node would fail to connect to a truncated (unexpected) address due to TLS peer verification. However, because this issue means that endpoint hosts longer than 255 bytes are broken, we would recommend disallowing setting hostnames or names longer than 255 bytes and validating that the appropriate sizes don't exceed the limits in the [UnpackNode](#) function.

2. The parseServerDuration implementation differs between [symbol-sdk-java](#) and [symbol-sdk-typescript-javascript](#). It seems that the TS/JS SDK implementation is missing a loop, so it can't parse complex durations such as "10h:10m", which are supported in the Java SDK. Additionally, the Java SDK implementation allows for "10h:10h" formats, which are most likely unexpected. We recommend:

- Fixing the parseServerDuration function in the TS/JS SDK implementation so it properly handles complex duration strings (such as "10h:10m").
- Fixing both implementations to disallow incorrect or unexpected inputs such as "10h:10h" or "10h:10m:10h".
- Making the tests between the two implementations consistent, so they check against similar cases.

- Changing the Java SDK function name from `parserServerDuration` to `parseServerDuration`, so that it matches the TS/JS naming.

3. Move the block type calculation to a separate function instead of performing the same calculation in multiple places. The following code paths calculate the block type:

- [catapult-server/plugins/coresystem/src/validators/BlockTypeValidator.cpp#L33](#)
- [catapult-server/extensions/harvesting/src/Harvester.cpp#L77](#)
- [catapult-server/tests/int/node/stress/test/BlockChainBuilder.cpp#L169](#)

4. While out of scope, the [TransactionStatusEnum](#) in the [symbol-openapi](#) project is missing some error codes. For example, the `Failure_Chain_*` errors are missing. We recommend investigating this issue and considering adding tests to validate this data or generating it from the source.

5. There is a typo "Coutn", instead of "Count" in Java SDK's [FinalizationStage enum](#).

H. Fix log for issues from Appendix G

NEM Group addressed issues listed in Appendix G in their codebase as a result of our assessment. Each of the fixes was verified by an audit team member.

1. The `endpoint.Host` and `metadata.Name` are truncated to a length of 255 (through the [GetPackedSize function](#) and a later `memcpy`) during Node objects serialization in the [PackNode function](#).

Fixed. The length of `endpoint.Host` and `metadata.Name` is now [validated when a Node object is created](#), ensuring the fields will not get truncated during serialization.

2. The `parseServerDuration` implementation differs between [symbol-sdk-java](#) and [symbol-sdk-typescript-javascript](#).

Fixed. The implementations ([symbol-sdk-java](#), [symbol-sdk-typescript-javascript](#)) should now accept the same duration formats.

3. Move the block type calculation to a separate function instead of performing the same calculation in multiple places.

Fixed. The code has been refactored and a [CalculateBlockTypeFromHeight function has been introduced](#).

4. While out of scope, the [TransactionStatusEnum](#) in the [symbol-openapi](#) project is missing some error codes.

Fixed. The error codes were added in [symbol-openapi#250](#).

5. There is a typo "Coutn", instead of "Count" in Java SDK's [FinalizationStage enum](#).

Fixed. The typo was fixed in a [recent commit](#).